

Accelerating I/O Performance for AI Frameworks on DAOS

Sooyoung Lim, Taehyun Yoo, Jaegi Son, and Dongmin Kim

slim@keti.re.kr, yootommy0113@gmail.com, jgson@keti.re.kr, dmkim@keti.re.kr


Korea Electronics Technology Institute (KETI)

Medical IT Convergence Research Center



Aurora-FS

Overview



As-is

❑ Heavy I/O overhead

- Modern AI and HPC workloads require intensive data preprocessing.


❑ Excessive redundant copies and context switches

- The default DAOS I/O path traverses between user space and kernel space via block drivers and RPC.

To-be


❑ Two solutions for accelerating I/O performance

- To reduce unnecessary data movement: **Zero-copy I/O**
- To move computation closer to data: **In-Storage Computing**




Motivation

- ❑ Setup – Evaluated PyTorch and TensorFlow preprocessing pipelines on DFUSE (DAOS FUSE) vs. POSIX filesystem (xfs) using image (CIFAR-10) and video (UCF101) datasets
- ❑ Result – Up to **3.3x slower** preprocessing on DFUSE for video pipelines
- ❑ Cause – DFUSE introduces frequent user-kernel transitions, redundant copies, and heavy POSIX operations
- ❑ Insight – Bottleneck stems from **data movement** between compute and storage!



Proposed Design



Redesigned DAOS I/O Path

(a) Zero-Copy I/O

- BIO → SPDK: Submits DMA-ready user buffers directly to SPDK's NVMe channels
- BIO ← SPDK: Returns completion events and signals the DMA transfer to devices

(b) In-Storage Computing

- User → BIO: Sends ISC requests to BIO by specifying the target object extents
- BIO → SPDK: Issues targeted block-access requests to SPDK

(a) Zero-Copy I/O

Limitation

- The software I/O path (user ↔ DFUSE ↔ kernel ↔ device) dominates latency.

Design Proposal

- Direct data transfer without kernel

Design Elements

❑ DAOS storage engines

- **BIO**: Acts as the zero-copy entry point that receives DMA-ready buffers from the DAOS client

❑ SPDK modules

- **Blobstore**: Manages physically contiguous I/O regions that can be directly mapped from user buffers
- **Bdev**: Provides DMA-capable I/O channels and NVMe queue pair management
- **SPDK memory module**: Handles hugepage-based physical address to make user buffers DMA-ready

(b) In-Storage Computing

Limitation

- Preprocessing workloads repeatedly transfer large volumes of I/O from compute to storage.

Design Proposal

- Offloading preprocessing pipelines to the storage layer

Design Elements

❑ DAOS storage engines

- **VOS**: Manages versioned object metadata and transactional consistency-guaranteed epochs

❑ BIO

- Hosts ISC execution and enables SPDK-driven I/O

❑ SPDK modules

- **Blobstore**: Executes ISC with block-granular access to NVMe-backed objects

- **Bdev**: Provides I/O channel control with block-level extent

❑ HW emulation

- **QEMU**: Provides a controllable NVMe device environment