

# Accelerating I/O Performance for AI Frameworks on DAOS




Sooyoung Lim, Taehyun Yoo, Jaegi Son, and Dongmin Kim

slim@keti.re.kr, yootommy0113@gmail.com, jgson@keti.re.kr, dmkim@keti.re.kr

Korea Electronics Technology Institute (KETI)

Medical IT Convergence Research Center

## Overview



### As-is

#### Heavy I/O overhead

- Modern AI and HPC workloads require intensive data preprocessing.


#### Excessive redundant copies and context switches

- The default DAOS I/O path traverses between user space and kernel space via block drivers and RPC.

### To-be


#### Two solutions for accelerating I/O performance

- To reduce unnecessary data movement: **Zero-copy I/O**
- To move computation closer to data: **In-Storage Computing**




## Motivation

- Setup** – Evaluated PyTorch and TensorFlow preprocessing pipelines on DFUSE (DAOS FUSE) vs. POSIX filesystem (xfs) using image (CIFAR-10) and video (UCF101) datasets
- Result** – Up to **3.3x slower** preprocessing on DFUSE for video pipelines
- Cause** – DFUSE introduces frequent user-kernel transitions, redundant copies, and heavy POSIX operations
- Insight** – Bottleneck stems from **data movement** between compute and storage!



## Proposed Design



### Redesigned DAOS I/O Path

#### (a) Zero-Copy I/O

- BIO → SPDK**: Submits DMA-ready user buffers directly to SPDK's NVMe channels
- BIO ← SPDK**: Returns completion events and signals the DMA transfer to devices

#### (b) In-Storage Computing

- User → BIO**: Sends ISC requests to BIO by specifying the target object extents
- BIO → SPDK**: Issues targeted block-access requests to SPDK

### (a) Zero-Copy I/O

#### Limitation

- The software I/O path (user ↔ DFUSE ↔ kernel ↔ device) dominates latency.

#### Design Proposal

- Direct data transfer without kernel

#### Design Elements

##### DAOS storage engines

- BIO**: Acts as the zero-copy entry point that receives DMA-ready buffers from the DAOS client

##### SPDK modules

- Blobstore**: Manages physically contiguous I/O regions that can be directly mapped from user buffers
- Bdev**: Provides DMA-capable I/O channels and NVMe queue pair management
- SPDK memory module**: Handles hugepage-based physical address to make user buffers DMA-ready

### (b) In-Storage Computing

#### Limitation

- Preprocessing workloads repeatedly transfer large volumes of I/O from compute to storage.

#### Design Proposal

- Offloading preprocessing pipelines to the storage layer

#### Design Elements

##### DAOS storage engines

- VOS**: Manages versioned object metadata and transactional consistency-guaranteed epochs

##### BIO

- Hosts ISC execution and enables SPDK-driven I/O

##### SPDK modules

- Blobstore**: Executes ISC with block-granular access to NVMe-backed objects

- Bdev**: Provides I/O channel control with block-level extent

##### HW emulation

- QEMU**: Provides a controllable NVMe device environment